

OPERATING SYSTEM TUTORIAL 2



PART-1

- start looping
- print: “username@hostname: /home/user >”
 - #include <unistd.h>
 - char *getlogin(void);
 - int gethostname(char *name, size_t namelen);
 - char *getcwd(char *buffer, size_t size);

Always check for errors, and free any memory that is allocated
Check manual for each function



University
of Victoria

PART-1

- read a line from terminal
- if `strcmp(line, "exit\n") == 0` (or `strcmp(line, "exit") == 10`), exit
- execute the input line by:
 - fork
 - execvp



execvp(char* file, char* argv[])

- #include <unistd.h>
- int execvp(const char *file, char *const argv[])
- execvp() provides **an array of pointers** to null-terminated strings that represent the argument list available to the new program.
- The array of pointers must be terminated by a NULL pointer. (tokenize input line)
- The exec family of functions **replaces** the current process image with a new process image.
- If execvp() returns -1, an **error** will have occurred.



Tokenize strings

- For calling `execvp(char* file, char* argv[])`, we need to separate input line by either empty space or new line (`'\n'`)

- example:

```
argv[0]=strtok(line," \n");//“ \n” includes space & new line
int i=0;
while(argv[i]!=NULL){
    argv[i+1]=strtok(NULL," \n");
    i++;
}
```



fork

- why need fork?
- example:

```
pid_t p=fork();  
if(p==0) {  
    execvp(argv[0], argv);  
    printf("ERROR\n");  
}
```

Negative Value: The creation of a child process was unsuccessful.

Zero: Returned to the newly created child process.

Positive value: Returned to parent or caller. The value contains the process ID of the newly created child process.



- parent needs to wait
- if a wait is not performed, resources associated with the child are not released. Then the terminated child remains in a “zombie” state



Wait

- `pid_t wait(int *status);`
`pid_t waitpid(pid_t pid, int *status, int options);`
- `pid`:
 - `<-1`: wait for any child process whose group id is `|pid|`
 - `-1`: wait for any child process
 - `0`: wait for any child process whose group id is equal to that of the calling process
 - `>0`: wait for the child process whose id is `pid`
- `options`:
 - `WUNTRACED` | `WCONTINUED` (or `0`): return if a child has stopped
 - `WNOHANG` (or `1`): return immediately if no child has exited

Wait

- return:
 - wait: on success, returns the process ID of the terminated child, return -1 if error
 - waitpid: on success, returns the process ID of the child whose state has changed. If WNOHANG was specified and child(ren) have not yet changed state, return 0. Return -1 if error.
- for example:
 - wait for any child until child terminates: wait(NULL) is shorthand for waitpid(-1, NULL, 0)
 - return immediately: waitpid(0, NULL, WNOHANG)
 - wait for certain child: waitpid(child_pid, NULL, 0)



PART-2

- if `strcmp(argv[0], "cd")==0` PART-2
else PART-1
- if `argv[1]==NULL || strcmp(argv[1], "~")==0`
go to HOME (env)
else go to `argv[1]`
- “echo \$HOME” in BASH to verify



PART-3

- `strcmp(argv[0], "bg") == 0`
- new argv: left shift the old argv, end with '\0'
- fork:
 - child: `execvp(newArgv[0], newArgv)`
 - parent:
 - if # of bg_process == 0
 - `root = bg_pro1`
 - else
 - append bg_proNext




linked list

- linked list for storing bg process info

```
struct bg_pro{  
    pid_t pid;  
    char command[1024];  
    struct bg_pro* next;  
};
```

- add, delete



bglist

- `strcmp(args[0], "bglist") == 0`
- print info in the linked list and number of bg processes



check if child terminates

- if # of bg_process > 0
pid_t ter=waitpid(0,NULL,WNOHANG);
 - if **ter > 0** //any child terminates
 - if root -> pid == ter
 - print root -> pid & root -> command & “has terminated”
 - root = root -> next
 - else
 - loop until we find cur -> pid == ter
 - print cur -> pid & cur -> command & “has terminated”
 - cur -> next = cur -> next -> next

Don't forget to free the removed nodes



University
of Victoria

Q&A

